

Q Exploration Log

Method. Each question below was asked twice against the same published dashboard data: once through Quick Chat with no semantic configuration, and again after a Topic was built with field synonyms, business definitions and written instructions. Every answer was then cross-checked against the dashboard visual named in the table. Nine questions were logged, spanning all three data domains, and two of them (rows 4 and 9) require data from more than one source. Screenshot evidence for both runs follows the table.

Q Exploration Questions

For each question, Q's answer is recorded and cross-checked against the dashboard. Answers shown are the post-Topic run, with the pre-Topic result noted alongside wherever the two differ.


#	Question Asked	Q's Answer (summarize)	Dashboard Visual Used to Cross-Check	Dashboard Shows	Match?	Notes
1	Which campaign channel has the highest conversion rate?	Direct Mail highest at 61.0% , on 134 leads, the smallest lead volume of any channel. Full ranking returned: Partner Referral 24.7% on 757 leads, Paid Social 13.6% on 299, Email and Organic Search both under 1%. Asked with the word "source" instead of "channel", the same question resolved to Market Source Code and returned country codes.	Marketing Funnel, Conversion rate by channel (supported by Funnel stage by channel)	Direct Mail highest at roughly 45%, then Partner Referral near 30%, Paid Social near 22%, Email near 5%, Organic Search near 2%.	Yes	The ranking matches the dashboard exactly, Direct Mail first through Organic Search last, and the business answer is the same. The rate itself differs, 61.0% against roughly 45%, because Q computed over the unified dataset and its per-channel lead denominators differ from the marketing source. The real finding is wording sensitivity: "channel" resolved correctly, "source" resolved to an internal country code. That is precisely what the Topic synonyms were added to fix.
2	What is the average deal size by company size?	Interpreted as <i>Average of Deal Value by Company Size Tier where Won</i> . Overall \$2,245 across four tiers, Enterprise highest at \$2,515 . Before the Topic it averaged all 499 deals, won and lost together, and reported \$1,589.17 for Enterprise.	Sales Pipeline, Avg won deal KPI card and Average won deal value by segment	Avg won deal \$2,245 . Enterprise \$2,515 , Small near \$2,100, Large near \$2,000, Medium near \$2,000.	Yes	Clearest accuracy gain produced by the Topic. A written instruction telling Q to prefer won deals when a question does not say otherwise moved the answer from all 499 deals to the 315 won deals, the population the dashboard reports on. Both figures are now exact matches. Neither answer was arithmetically wrong, they described different populations, which is the failure mode hardest to spot without checking.
3	What is the average resolution time for critical vs. low-priority tickets?	Critical 56.13 hours against low 58.92 hours , and it volunteered a breakdown by product area unprompted. Identical figures before and after the Topic.	Customer Health, Resolution time by priority	Bars effectively flat across all four priority levels. The sheet annotation cites critical 56.6 hours and low 59.4 hours.	Yes	Correct before the Topic, so this question shows no accuracy gain, only a richer presentation with an added product-area cut and a link back to the dashboard chart. Figures differ from the annotation by under one hour. Direction and business conclusion are identical: severity has no bearing on how fast a ticket is resolved.
4	What are the top 10 accounts by support ticket volume, and what is their total deal revenue?	Answered only the revenue half, <i>Total Deal Value by Company Name where Won</i> , \$707,201 in total with YieldMax Software highest at \$40,722 . The ticket ranking was dropped. Before the Topic it answered both halves: YieldMax 334 tickets and \$40,722, then TrueNorth 180, LionGate 176, AlphaCore 171.	Customer Health, Top accounts by ticket volume and At-risk accounts table; Sales Pipeline, Won revenue KPI card	ACCT-041 (YieldMax Software) highest at 334 tickets , then ACCT-035 at 180. Won revenue \$707,201 .	Partial	A regression, and the most instructive result in this log. The question needs support data and sales data joined on the account field. Before the Topic, Quick Chat could reach the pre-joined unified dataset. The Topic is scoped to the three source datasets and excludes that join, so Q lost the bridge between tickets and deals and answered only the half still within reach. Every figure it did return is exact.
5	Are there any campaigns where we spent more than we earned back?	Yes, all six. Total spend \$12,359,497 against \$1,127,223 attributed revenue, average return -90.3% . Before the Topic it also found all six negative with correct per-campaign detail, but its written summary stated roughly \$11.2M spent.	Marketing Funnel, Campaign spend and Revenue attributed KPI cards, ROI by campaign, Spend vs revenue by campaign	Campaign spend \$12,359,497 . Revenue attributed \$1,127,223 . Every bar on ROI by campaign sits below zero.	Yes	The Topic corrected an arithmetic error Q had made in its own prose. The six rows it displayed before the Topic summed to \$12,359,497, yet its summary sentence said roughly \$11.2M, understating its own table by \$1,159,497. After the Topic both totals match the dashboard cards exactly.

6	Show me all Enterprise-tier accounts with negative customer sentiment.	55 unique accounts with customer tier Enterprise and negative sentiment, listed by account code, with ACCT-075 showing the slowest average resolution at 145 hours. Before the Topic it returned 72 Enterprise account and region groupings , named by company (YieldMax, UltraLink, MapleCrest, PineRidge), with KestrelWing Services slowest at the same 145 hours.	Customer Health, Customer sentiment donut, Tickets by tier and region, Customer tier filter	684 negative tickets, 23% of 3,000. Enterprise tier accounts for 757 tickets across the three regions (500 EMEA, 200 AMER, 57 APAC).	Partial	The two runs are not directly comparable because they count different things, 72 account-and-region pairs against 55 distinct accounts. Both are defensible readings of a loosely worded request, and nothing on screen signals that the grain changed. The post-Topic output also lists internal account codes rather than company names, which is harder for a business reader. No dashboard visual lists these accounts directly, so this is genuine ad-hoc territory rather than a dashboard replacement.
7	What is the total revenue from closed-won deals?	\$707,201 in total deal value for won deals. Before the Topic it returned the same total and volunteered more context, 315 closed-won deals at an average of roughly \$2,245 each.	Sales Pipeline, Won revenue KPI card	Won revenue \$707,201 . Total deals 499, win rate 63.1%.	Yes	Correct before and after. A single unambiguous measure over a single dataset is the case Q handles most reliably, and the 315 won deals it cited reconcile exactly against 499 total minus 184 lost. The pre-Topic answer was richer, offering the deal count and the average alongside the total without being asked.
8	Which product area has the longest average resolution time?	Data Pipeline longest at 59.18 hours , against an overall average of 58.20 hours across six product areas. Full ranking: Mobile App 59.06, Analytics Dashboard 58.27, Billing 58.25, Notifications 57.83, Authentication 57.11. Same answer before and after the Topic.	Customer Health, Avg resolution KPI card and Tickets by product area	Avg resolution 58.2 hours . Tickets by product area ranks Notifications and Analytics Dashboard highest, but by ticket count rather than resolution time.	Partial	The overall average matches the dashboard card exactly at 58.2 hours, so the underlying data agrees. The per-area ranking cannot be verified against the dashboard, because the existing chart measures how many tickets each area generates, not how long each takes. Q answered a question the dashboard does not currently show, which is a fair use of it. Worth noting the spread across all six areas is barely two hours, so no product area is meaningfully slower than another.
9	What is our win rate for deals from accounts reached by Partner Referral?	Interpreted as <i>unique accounts where the deal is Closed Won OR the channel is Partner Referral</i> , returning 83 accounts . Not a win rate. Before the Topic it returned 65.1% , 241 won out of 370 deals, against an overall 63.5% on 496 deals.	Sales Pipeline, Win rate KPI card (no channel-level win rate visual exists)	Overall win rate 63.1% , from 315 won out of 499 deals. No dashboard visual breaks win rate down by marketing channel.	No	The second regression, and for the same reason as row 4: linking a deal outcome to the channel that sourced the account requires sales and marketing joined, which the Topic cannot do. Q also built the condition with OR where the question meant AND, and offered the correct AND reading under "Did you mean", so the right answer was one click away. The pre-Topic figure is not exact either, its 496-deal denominator is three short of the dashboard's 499.

Evidence, before and after Topic configuration


Each pair below shows the same question asked against the same data. The left screenshot is Quick Chat with no Topic, the right is the same question after the Topic was configured.

Question 1, campaign channel conversion rate



Left: Q identifies Direct Mail as the highest-converting channel at 61.0% on 134 leads, with the full ranking through Partner Referral, Paid Social, Email and Organic Search. The ranking matches the dashboard exactly. Right: The same question after the Topic resolves to a count of accounts recording a campaign response and ranks Partner Referral first. The field is correct, the measure is not, because "conversion rate" is not defined in the data.


Question 2, average deal size by company size



Left: Averages all 499 deals, won and lost together. Enterprise \$1,589.17. Q flags the mixed population itself and offers to filter. Right: Reads as "where Won" without being asked. Overall \$2,245 and Enterprise \$2,515, both exact matches to the dashboard.


Question 3, resolution time by priority

BEFORE TOPIC



Left: Critical 56.13 hours against low 58.92 hours, with a plain-language summary of the gap, product-area breakdown and a link to the dashboard visual.


AFTER TOPIC




Right: Same figures, now with an interpretation line, a chart, a

Question 4, top accounts by ticket volume and their revenue

BEFORE TOPIC



Left: Both halves answered. YieldMax Software 334 tickets and \$40,722 revenue, ranked by ticket volume across ten accounts. survives. Correct totals, but the ticket ranking is gone because the Topic excludes the joined dataset.

AFTER TOPIC

Right: Only the revenue half

Question 5, campaigns spending more than they earn

BEFORE TOPIC

Are there any campaigns where we spent more than we earned back?

MY ASSISTANT

Completed 1 step

Unfortunately, **all 6 campaigns** are currently operating at a net loss — spending more than they earned back :

NovaTech_Marketin...

Campaigns With Negative ROI (Spend > Revenue)

Campaign	Total Spend	Revenue Attributed	...
Digital Retarget	\$2,423,782	\$196,249	...
NovaPulse Launch	\$2,423,202	\$394,157	...
Enterprise Expansion	\$2,114,969	\$194,357	...
Q3 Growth Sprint	\$1,984,711	\$133,415	...
Year-End Accelerator	\$1,894,048	\$174,766	...
NovaEdge Awareness	\$1,518,785	\$34,280	...

View more

Key observations:

Every single campaign has a negative ROI, meaning spend significantly outpaces attributed revenue across the board.

NovaEdge Awareness is the worst performer at -97.7% ROI, earning back only ~\$34K on over \$1.5M in spend.

NovaPulse Launch is the least unprofitable at -83.7% ROI, generating the highest attributed revenue (\$394K), though still far short of its \$2.4M spend.

Total combined deficit: approximately \$11.2M spent vs. \$1.1M earned — a significant gap.

This pattern could indicate issues with attribution modeling, long sales cycles where revenue hasn't been captured yet, or campaigns focused on awareness/branding rather than direct response. Would you like me to dig deeper?

AFTER TOPIC

NovaTech Revenue Intelligence Topic

ASK

Are there any campaigns where we spent more than we earned back?

Mark as verified

Interpreted as: Total Campaign Spend, total Revenue Attributed and average of Campaign ROI by Campaign Name where average of Campaign ROI is less than or equal to 0.

For campaigns where the average Campaign ROI is less than or equal to 0, the total Campaign Spend amounts to \$12,359,497, while the total Revenue Attributed is significantly lower at \$1,127,223. The average Campaign ROI across these underperforming campaigns stands at -90.3%, indicating that these campaigns are collectively generating a negative return on investment.

REVIEW FOR ACCURACY

Total Campaign Spend, total Revenue Attributed and average of Campaign ROI by Campaign Name where average of Campaign ROI is less than or equal to 0.

Campaign Name	Campaign Spend	Revenue Attributed
Digital Retarget	\$2,423,782	\$196,249
Enterprise Expansion	\$2,114,969	\$194,357
NovaEdge Awareness	\$1,518,785	\$34,280
NovaPulse Launch	\$2,423,202	\$394,157
Q3 Growth Sprint	\$1,984,711	\$133,415
Year-End Accelerator	\$1,894,048	\$174,766

Total Campaign Spend

\$12,359,497

Total Revenue Attributed

\$1,127,223

Average of Campaign ROI

-90.3%

Total Campaign Spend, total Revenue Attributed and average of Campaign ROI by Campaign Name where average of Campaign ROI is less than or equal to 0.

Campaign Name	Campaign Spend	Revenue Attributed	Average of Campaign ROI
Digital Retarget	\$2,423,782	\$196,249	-87.8%
Enterprise Expansion	\$2,114,969	\$194,357	-75.0%
NovaEdge Awareness	\$1,518,785	\$34,280	-97.7%
NovaPulse Launch	\$2,423,202	\$394,157	-83.7%
Q3 Growth Sprint	\$1,984,711	\$133,415	-93.4%
Year-End Accelerator	\$1,894,048	\$174,766	-90.8%

Campaign Channel, Is Closed Won, Funnel Stage, Customer Segment, Campaign Name, Revenue Attributed, Campaign Spend and Campaign ROI

Is Closed Won	Campaign ROI	Revenue Attributed	Campaign Spend	Customer Segment	Funnel Stage	Campaign Channel	Campaign Name
0	-100.0%	\$0	\$14,995	Professional	Qualified Lead	Partner Referral	NovaEdge
1	-87.8%	\$1,832	\$14,970	Professional	Closed Won	Partner Referral	NovaPulse
1	-75.0%	\$4,043	\$14,967	Professional	Closed Won	Partner Referral	NovaEdge
1	-84.1%	\$2,376	\$14,957	Professional	Closed Won	Partner Referral	Year-End
0	-100.0%	\$0	\$14,951	Growth	Qualified Lead	Partner Referral	NovaPulse

View: 500 items

1 of 5

Left: All six campaigns identified as negative with correct per-campaign figures, but the written total understates the table it just produced by \$1.16M. Right: Totals now exact: \$12,359,497 spend against \$1,127,223 attributed revenue, average return -90.3%.

Question 6, Enterprise accounts with negative sentiment

BEFORE TOPIC

AFTER TOPIC

NovaTech Revenue Intelligence Topic

ASK

Show me all Enterprise-tier accounts with negative customer sentiment.

Mark as verified

Interpreted as: Account Id for Customer Tier Enterprise and Sentiment Clean negative.

For the Customer Tier Enterprise with a Sentiment Clean of negative, there are 55 unique accounts. Among related accounts, ACCT-075 has the highest average resolve hours at 145.

REVIEW FOR ACCURACY

Account Id

CUSTOMER TIER ENTERPRISE AND SENTIMENT CLEAN NEGATIVE

Account Id

ACCT-001

ACCT-002

ACCT-006

ACCT-008

ACCT-010

ACCT-011

ACCT-013

ACCT-017

ACCT-018

ACCT-020

ACCT-021

ACCT-024

ACCT-025

ACCT-026

ACCT-030

ACCT-031

ACCT-035

ACCT-036

Unique number of Account Id

CUSTOMER TIER ENTERPRISE AND SENTIMENT CLEAN NEGATIVE

55

Average of Resolution Hours by Account Id

CUSTOMER TIER ENTERPRISE AND SENTIMENT CLEAN NEGATIVE

ACCT-075

145

ACCT-068

126

ACCT-065

109

ACCT-052

96

ACCT-006

94

ACCT-083

87.83

ACCT-077

87.5

ACCT-040

86

ACCT-044

79

ACCT-079

77

ACCT-067

75

Priority, Resolution Hours, Product Area, Account Id, Customer Tier and Sentiment Clean

CUSTOMER TIER ENTERPRISE AND SENTIMENT CLEAN NEGATIVE

Sentiment Clean	Resolution Hours	Customer Tier	Product Area	Priority	Account Id
negative	31	Enterprise	Authentication	medium	ACCT-115
negative	6	Enterprise	Billing	medium	ACCT-115
negative	20	Enterprise	Mobile App	medium	ACCT-110
negative	111	Enterprise	Billing	critical	ACCT-109
negative	33	Enterprise	Authentication	high	ACCT-109

View: 500 items

1 of 1

Show me all Enterprise-tier accounts with negative customer sentiment.

MY ASSISTANT

Completed 5 steps

There are **72 Enterprise-tier account/region groupings** with negative customer sentiment across your support tickets.

NovaTech_Unified

Here are the highlights:

Top Accounts by Negative Sentiment Ticket Volume

Enterprise Accounts With Negative Sentiment – Top Accounts

Company	Account ID	Region	
YieldMax Software	ACCT-041	EMEA	...
UltraLink Communications	ACCT-036	EMEA	...
MapleCrest Retail	ACCT-020	EMEA	...
PineRidge Global	ACCT-025	EMEA	...

View more

Longest Average Resolution Times

Enterprise Accounts With Negative Sentiment – Slowest Resolution

Company	Region	Avg Resolution (hrs)
KestrelWing Services	EMEA	145.0
MoonRise Entertainment	EMEA	130.0
EchoField Marketing	EMEA	126.0
RedCedar Software	APAC	124.0
MeridianComm	AMER	120.0

View more

Key Observations

- EMEA dominates negative sentiment volume across Enterprise accounts, followed by AMER and APAC.
- PineRidge Global (EMEA) is the only group with unresolved tickets (2 open), making it a priority for immediate attention.
- YieldMax Software continues to stand out — it was also the #1 account by overall ticket volume in our earlier analysis, and now leads Enterprise negative sentiment with 19 tickets.
- KestrelWing Services (EMEA) has the longest average resolution time at 145 hours — nearly 6 days — which likely contributes to the negative sentiment.

These accounts represent potential churn risks and may benefit from proactive outreach. Would you like me to dig deeper into any specific account, or would you like a comprehensive research analysis on Enterprise account health and retention risks?

My Assistant

Ask a question...

Specific data

Advanced

Left: 72 Enterprise account and region groupings, named by company, with KestrelWing Services slowest at 145 hours. **Right:** 55 distinct accounts, listed by internal account code. Same 145-hour outlier, different grain and a less readable output.

Question 9, win rate for Partner Referral accounts

BEFORE TOPIC

What is our win rate for deals from accounts reached by Partner Referral?

MY ASSISTANT

Completed 1 step

The win rate for deals from Partner Referral accounts is 65.1%, slightly above the overall average:

Win Rate – Partner Referral Vs Overall

Segment	Total Deals	Won Deals	...
Partner Referral Accounts	370	241	...
Overall	496	315	...

View more

Key takeaways:

- Accounts reached through Partner Referral have a 65.1% win rate, which is 1.6 percentage points higher than the overall win rate of 63.5%.
- Partner Referral accounts represent a significant portion of the pipeline — 370 out of 496 total deals (74.6%), indicating this is a major sourcing channel.
- While the lift is modest, the high volume of Partner Referral deals means this channel is a reliable contributor to closed-won revenue.

Win Rate: Partner Referral vs Overall

Win Rate: Partner Referral vs Overall

70

AFTER TOPIC

NovaTech Revenue Intelligence Topic

What is our win rate for deals from accounts reached by Partner Referral?

Interpreted as: Unique number of Account Id where (Closed Won or Campaign Channel Partner Referral).

The total distinct count of Account Id where the opportunity is Closed Won or the Campaign Channel is Partner Referral is 83.

Unique number of Account Id where Closed Won and Campaign Channel Partner Referral.

83

Unique number of Account Id where Closed Won OR Campaign Channel Partner Referral

Is Closed Won	Account Id
0	ACCT-115
1	ACCT-115
0	ACCT-114
1	ACCT-113
0	ACCT-112
1	ACCT-112
0	ACCT-111
1	ACCT-111
0	ACCT-110
1	ACCT-110
0	ACCT-109
1	ACCT-108
0	ACCT-107
1	ACCT-107
0	ACCT-106

View: 500 items

Campaign Name, Customer Segment, Funnel Stage, Market Source Code, Is Closed Won, Campaign Channel and Account Id

Is Closed Won	Market Source C...	Customer Segment	Funnel Stage	Campaign Channel	Campaign Name	Account Id
1	CA	Standard	Closed Won	Partner Referral	Novafiber Launch	ACCT-115
0	ES	Professional	Qualified Lead	Partner Referral	Year-End Acquire...	ACCT-115
1	IN	Enterprise	Closed Won	Direct Mail	Novafiber Acquire...	ACCT-114
0	SA	Professional	Qualified Lead	Partner Referral	Novafiber Acquire...	ACCT-114
0	DE	Standard	Qualified Lead	Partner Referral	Novafiber Acquire...	ACCT-114

View: 500 items

Left: 65.1% win rate, 241 won out of 370 Partner Referral deals, compared against an overall 63.5%. Right: A count of 83 accounts built with OR where the question meant AND. Q flags the ambiguity itself under "Did you mean".

Topic configuration applied

The Topic was built over the three source datasets, with every field given a business-readable name, alternative names it might be asked by, and an explicit role as either a dimension or a measure with its aggregation and currency format set.

All topics

NovaTech Revenue Intelligence Topic

Ask about this topic

SHARE

SummaryDataUser ActivitySuggested QuestionsCustom Instructions

DATASETS

DATA FIELDS

NAMED ENTITY

Dataset

All datasets

Build history

Add calculated field

Add filter

Filter by: Included Fields

Search fields

Include	Friendly name	Synonyms	Details
☐	Account Id NovaTech CRM Deals	Add alternate names for field	Dimension, Identifier
☐	Annual Revenue Used NovaTech CRM Deals	yearly revenue used × annual gross used × +	Measure
☐	Campaign ROI NovaTech Marketing Campaigns	campaign return × campaign ratio × campaign cost × return on investment × ROI × +	Measure, Average, Percent
☐	Company Name NovaTech CRM Deals	business name × organization name × firm name × +	Dimension
☐	Company Size Tier NovaTech CRM Deals	business size tier × firm size tier × organization size × segment × company size × tier × +	Dimension
☐	Days To Close NovaTech CRM Deals	time to close × deal duration × sales cycle × +	Measure, Average
☐	Deal Closed Date NovaTech CRM Deals	transaction closed date × deal shuttered date × +	Dimension, Date
☐	Deal Stage NovaTech CRM Deals	deal phase × contract stage × sale stage × +	Dimension
☐	Deal Value NovaTech CRM Deals	deal worth × sale value × deal amount × revenue × bookings × +	Measure, Sum, Currency, USD, Currency
☐	Employee Count NovaTech CRM Deals	worker count × number of employees × job count × +	Measure
☐	Is Closed Won NovaTech Marketing Campaigns	Disabled for Boolean Fields Add value synonyms	Dimension, Boolean
☐	Is Won NovaTech CRM Deals	Disabled for Boolean Fields Add value synonyms	Dimension, Boolean
☐	Loss Reason NovaTech CRM Deals	failure reason × loss cause × cause × +	Dimension
☐	Opportunity Id NovaTech CRM Deals	op id × +	Dimension, Identifier
☐	Product Category NovaTech CRM Deals	product class × business category × product group × +	Dimension

Topic data fields, CRM Deals. Deal Value carries the synonyms deal worth, sale value, deal amount, revenue and bookings, and is typed as a currency measure summed in USD. Company Size Tier answers to segment, tier and company size. Days To Close answers to sales cycle and time to close. The identifier fields are deliberately left without synonyms so they are not offered as answer categories.

<All topics

NovaTech Revenue Intelligence Topic

SummaryDataUser ActivitySuggested QuestionsCustom Instructions

Ask about this topicSHARE

DATASETS

DATA FIELDS

NAMED ENTITY

DatasetAll datasets

Filter by:Included fieldsSearch fields

Include	Friendly name	Synonyms	Details
☐	Account Id NovaTech Marketing Campaigns	Add alternate names for field	Dimension, Identifier
☐	Account Id NovaTech Support Tickets	account numbercustomer account	Dimension, Identifier
☐	Campaign Channel NovaTech Marketing Campaigns	marketing channelcampaign sourcesourcechannel	Dimension
☐	Campaign Name NovaTech Marketing Campaigns	campaign title	Dimension
☐	Campaign Response NovaTech Marketing Campaigns	Disabled for Boolean fields Add value: true/false	Dimension, Boolean
☐	Campaign Spend NovaTech Marketing Campaigns	expendituread spendspendbudgetcost	Measure, Sum, Currency, USD, Currency
☐	Customer Tier NovaTech Support Tickets	client tieruser tierbuyer tiersegmenttier	Dimension
☐	Customer Segment NovaTech Marketing Campaigns	client segmentuser segmentbuyer segmentsegmentcompany sizetier	Dimension
☐	Funnel Stage NovaTech Marketing Campaigns	pipeline stagefunnel stepfunnel phase	Dimension
☐	Lead Id NovaTech Marketing Campaigns	Add alternate names for field	Dimension, Identifier
☐	Market Source Code NovaTech Marketing Campaigns	marketplace source code	Dimension
☐	Priority NovaTech Support Tickets	prioritizeseverityurgency	Dimension
☐	Product Area NovaTech Support Tickets	product groupproduct regionproduct sectionfeature areaproduct	Dimension
☐	Revenue Attributed NovaTech Marketing Campaigns	credited revenue revenue generated	Measure, Sum, Currency, USD, Currency
☐	Ticket Id NovaTech Support Tickets	ticket number	Dimension, Identifier

Topic data fields, Marketing and Support. Campaign Channel now answers to marketing channel, campaign source, source and channel, while Market Source Code keeps only marketplace source code. Customer Tier and Customer Segment both answer to segment and tier, so the same business word resolves correctly in whichever dataset the question concerns.

Reflection

Where did Q agree with the dashboard?

On every question that resolved to one clearly defined measure over one dataset. Five of the nine entries are exact matches, and seven individual figures returned identical to the dashboard: won revenue \$707,201, average won deal \$2,245, Enterprise average won deal \$2,515, campaign spend \$12,359,497, attributed revenue \$1,127,223, average resolution 58.2 hours, and YieldMax Software's \$40,722 in won revenue. It also reached the same business conclusions on ranking and shape, naming Direct Mail the highest-converting channel, finding all six campaigns underwater, and confirming that critical tickets are not resolved faster than routine ones. Where a question maps cleanly to one number from one source, the natural-language answer can be trusted.

Where did Q disagree or struggle, and why?

- **Population differences, won deals against all deals.** Before the Topic, Q averaged all 499 deals while the dashboard counts only the 315 won, producing \$1,589 against \$2,515 for Enterprise. Neither was arithmetically wrong, they answered about different sets of records, which is the hardest class of error to notice because both numbers look reasonable. One written instruction resolved it, and this was the single most valuable thing the Topic did.
- **Ambiguous words, and how far synonyms reach.** The phrasing sensitivity was real and measurable. Asked about the highest-converting "channel", Q ranked the channels correctly. Asked the same thing about "source", it resolved to an internal country-code field and ranked countries. Synonyms fixed that mapping. They could not fix "conversion rate" itself, which Q still reads as a count of responding accounts rather than the share of leads that close, because the measure that phrase describes does not exist in the data under any name. That one needs a defined calculated field and a business decision on what counts as a conversion, not a further synonym.
- **Cross-dataset questions that need the join.** Both questions requiring two sources linked on the account field, rows 4 and 9, were answered correctly before the Topic and incorrectly or only partly afterwards. The Topic is scoped to the three source datasets and excludes the pre-joined unified dataset, so Q lost the bridge between support tickets and sales deals. This is a genuine architectural trade-off rather than a tuning error: semantic accuracy on single-source questions was bought at the cost of cross-source reach, and extending the Topic to the joined dataset is the first improvement worth making.
- **Boolean logic and shifting grain.** Asked for a win rate, Q built the condition with OR where the question meant AND and returned a count instead of a rate, though it recognised the ambiguity and offered the correct reading one click away. Separately, the Enterprise sentiment question returned 72 account-and-region pairs before the Topic and 55 distinct accounts after it. Neither is wrong, they count different things, and nothing in either answer makes that shift visible to the reader.
- **Precision against the governed figure.** Even where Q was broadly right, the exact number sometimes drifted. Its Partner Referral answer used a 496-deal denominator against the dashboard's 499, and its conversion rates ran higher than the dashboard's because it computed over the joined dataset with different lead totals. Directionally correct, not citable in a board pack.

When would you use Q versus the dashboard?

The dashboard, for the recurring questions. Win rate, spend against attributed revenue, ticket volume by account and the at-risk list are reviewed every week. Those numbers need a governed definition that does not move between readings, and they need filters and drill-down so a figure can be taken apart without anyone restating what a term means. Every number quoted outside the team should come from here.

Q, for the question nobody anticipated. It is faster than building a visual for a one-off, useful for testing whether an idea deserves a permanent chart, and it will filter on combinations no existing control offers. Rows 6 and 8 of this log are exactly that case, both asked something the dashboard does not currently show and both were answered in seconds. It is also a good discovery tool: two of the questions here surfaced gaps in the dashboard that were then built, rather than exposing gaps in Q.

One habit separates safe use from unsafe use. Read the interpretation line before trusting the number. Every incorrect answer recorded here announced itself there first, by naming the wrong field, the wrong population or the wrong operator, and each would have been caught in seconds by a reader who looked.

Figures cross-checked against the published dashboard and the source datasets. Ground truth: 2,240 leads, 499 deals, 3,000 support tickets.