

Executive Summary: NovaTech Revenue Performance

Marketing, sales and customer health across 85 accounts

Prepared for: Sarah Chen, VP of Revenue | **Source:** NovaTech Revenue Intelligence Dashboard, Amazon Quick with Amazon Q | **Period:** Deals closed December 2023 to January 2025

This summary covers \$707,201 of closed-won revenue from 499 sales opportunities, \$12,359,497 of marketing spend across six campaigns and 2,240 leads, and 3,000 support tickets across 85 accounts. It sets out the six findings most relevant to the revenue team, each verified against both the dashboard and an independent natural-language query.

Executive Snapshot

- Every one of the six active campaigns returned less than it cost, \$12.36M spent against \$1.13M attributed.
- Support resolution time ignores priority entirely, critical tickets wait as long as routine requests.
- Eight accounts combine high ticket volume, negative sentiment and above-median revenue, and nobody owns that signal today.
- The highest-revenue account in the business is also the highest support burden, at 334 tickets.
- 47% of lost deals were lost to no decision or poor fit, not to price or a competitor.

Key Insights

1. Marketing Spend Is Not Returning Revenue

Observation: All six active campaigns spent more than they earned back. \$12,359,497 in campaign spend returned \$1,127,223 in attributed revenue, an average return of -90.3%. NovaEdge Awareness is the weakest at -97.7%, recovering roughly \$34,000 on over \$1.5M of spend. NovaPulse Launch is the strongest and is still at -83.7%.

Meaning: A shortfall this uniform across six independently run campaigns is more likely to be one measurement problem than six separate failures. Account-level tracing confirms it: Falcon Software received \$18,473 of campaign spend, is recorded in marketing as having converted nothing, and attributes \$0 of revenue, yet sales closed nine deals for that same account worth \$20,912. The revenue exists and marketing is not crediting it. Cutting budget on these numbers risks cutting campaigns that are working.

Action: Head of Marketing with the Data and Analytics Lead, audit the attribution rules and confirm the revenue window before any budget decision. Two weeks.

2. The Channels That Convert Best Receive the Least Volume

Observation: Direct Mail converts leads to closed deals at roughly 45% and Partner Referral at roughly 30%. Organic Search converts near 2% and Email near 5%, yet Organic Search carries the largest share of total leads of any channel.

Meaning: The lead mix is close to inverted against performance. The two channels that reliably produce revenue are being fed the least, while the highest-volume channel is the weakest converter in the portfolio.

Action: Head of Marketing, present a reallocation plan shifting spend toward Direct Mail and Partner Referral, with a volume ceiling stated for each. Next monthly review.

3. A High Win Rate Is Concealing Small Deals

Observation: Central wins 69.9% of its deals at an average of \$1,905 per won deal. West wins 58.8% but averages \$2,559, and East averages \$2,496. Across the business the win rate is 63.1% on 499 closed deals worth \$707,201.

Meaning: On the measure currently reported, Central is the strongest region. On revenue per won deal it is the weakest. A target built on win rate alone rewards pursuing small, easy deals and quietly penalises the regions bringing in more money per sale.

Action: VP of Revenue with the six Sales Managers, add revenue per won deal to the regional scorecard and review Central's deal mix. Next quarterly planning cycle.

4. Nearly Half of All Losses Were Avoidable

Observation: Of 184 lost deals, No Decision Made accounts for 43 and Poor Product Fit for another 43. Together they are 86 losses, 47% of the total. Neither is about price and neither involves a competitor. The average deal takes 66.8 days to close.

Meaning: These are qualification failures rather than competitive losses. No decision means the buyer never committed, poor fit means the deal should not have been pursued. Both consumed a full sales cycle each, which puts roughly 5,700 selling days into deals that were never winnable.

Action: Sales Managers, introduce a qualification gate, budget confirmed, decision maker identified and product fit verified, before any deal advances to Opportunity stage. Next sprint.

5. Urgent Support Tickets Are Not Treated as Urgent

Observation: Resolution time is effectively flat across every priority level. Critical tickets average 56.6 hours and low-priority tickets 59.4 hours, a difference under three hours. Across the six product areas the spread is under two hours. Fifty tickets are classified critical, including 14 security incidents and 11 involving data loss.

Meaning: Priority is recorded but has no operational effect. A security incident currently waits about as long as a routine request. The classification exists in the data and is not reflected in how work is scheduled.

Action: Support Lead, define priority-based service levels starting with an eight-hour target for critical tickets, and report attainment weekly. Thirty days.

6. Eight Accounts Are at Real Risk of Churning Quietly

Observation: Eight accounts combine above-median ticket volume, 25% or more negative sentiment and above-median revenue. Falcon Software is the clearest case, \$20,912 in won revenue with 41% of its tickets filed in a negative tone. Separately, YieldMax Software files 334 tickets, nearly double the next account, while holding the highest won revenue in the business at \$40,722.

Meaning: These accounts are paying customers showing sustained dissatisfaction, and nobody currently owns that signal. The largest account by revenue is also the largest by support load, which is the profile that churns without warning.

Action: VP of Revenue, run executive account reviews across the eight named accounts and assign dedicated support coverage to YieldMax Software. This quarter.

Confidence Summary

Insight	Confidence	Verification method
1. Marketing spend against return	High	Campaign spend and Revenue attributed KPI cards against an Amazon Q total-spend query, both returning \$12,359,497 and \$1,127,223. The attribution hypothesis was then tested at account level by filtering all three views to a single account.
2. Channel conversion performance	Medium	Conversion rate by channel visual. Amazon Q was asked the same question and resolved 'conversion' to a different measure, so the dashboard figure is the governed one.
3. Win rate against deal size	High	Win rate by region and Average won deal value by region, cross-checked against the 499-deal total and the \$707,201 won revenue card.
4. Loss reason concentration	High	Loss reasons visual against the 184 lost deals derived from the 499 total and 315 won.
5. Resolution time by priority	High	Resolution time by priority visual against an Amazon Q query returning 56.13 and 58.92 hours, and against the 58.2 hour average card.
6. At-risk accounts	High	At-risk accounts table built on the joined dataset, cross-checked account by account against ticket volume, sentiment and won revenue.

Every figure was re-queried in Amazon Q and matched the dashboard once both views were set to the full unfiltered period. One discrepancy remains open and is recorded above: the term "conversion rate" resolves to different measures in the two systems, because it is not defined as a field in the underlying data. The dashboard figure is authoritative until that definition is agreed.

Emerging Risks

- Attribution integrity.** A -90.3% average return across all six campaigns undermines confidence in every marketing figure until the cause is identified. At least one account earns real revenue that marketing records as \$0, so the gap is demonstrable rather than theoretical.
- Undetected churn.** Eight accounts carrying real revenue are showing sustained dissatisfaction with no owner and no review cadence.
- Single-account concentration in support.** One account generates 11% of all support tickets while also holding the largest revenue balance.
- Unenforced severity.** Fourteen security incidents and eleven data-loss tickets are currently handled on the same timeline as routine requests.
- Definition drift.** Conversion rate, customer size tier and won-versus-all-deals each mean different things in different systems, which will produce conflicting numbers in meetings until they are standardised.

Recommended Next Steps

- Audit campaign attribution before any budget decision, Head of Marketing with Data and Analytics Lead, two weeks.
- Reallocate lead spend toward Direct Mail and Partner Referral, Head of Marketing, next monthly review.
- Add revenue per won deal to the regional scorecard, VP of Revenue with Sales Managers, next quarterly cycle.
- Introduce a qualification gate before Opportunity stage, Sales Managers, next sprint.
- Define priority-based service levels starting at eight hours for critical, Support Lead, thirty days.
- Run executive reviews on the eight at-risk accounts and assign dedicated coverage to YieldMax Software, VP of Revenue, this quarter.
- Agree a single definition of conversion rate and publish it to the dashboard and the query layer, Data and Analytics Lead, before the next reporting cycle.

Data sources: NovaTech Revenue Intelligence Dashboard (Marketing Funnel, Sales Pipeline, Customer Health) · Amazon Q Topic analysis · NovaTech CRM, marketing and support datasets, 2023 to 2025.