

# NovaTech Revenue Intelligence Dashboard

Findings and recommendations for the VP of Revenue

---

**Prepared for:** Sarah Chen, VP of Revenue

**Summary.** The dashboard you asked for is live. It brings marketing, sales and support into one place, three views instead of three systems, and the same account can now be followed from the campaign that generated the lead, through the deal, to the support tickets that came after. Seven findings are set out below. The two that need attention first are the eight accounts at risk of churning quietly, and the fact that urgent support tickets are currently resolved no faster than routine ones.

## How the data was put together

---

Three sources feed the dashboard: 499 sales opportunities from the CRM, 2,240 marketing leads across six campaigns, and 3,000 support tickets. They were loaded into SPICE, the in-memory store that Amazon Quick uses to keep a dashboard responsive, and refreshed on a weekly schedule, which matches the freshness you said you needed.

Each source was checked against the supplied data dictionary before anything was built. That check caught two errors worth knowing about. Opportunity identifiers are not unique on their own, three are reused across different accounts, so a deal is only safely identified by account and opportunity together. Separately, the support system records three customer sizes while the CRM records four, so the two systems do not label a customer's size identically.

The three sources were then joined into one combined dataset using the account identifier, the one field that appears in all three. Sales is the anchor, with marketing and support attached to it. That combined view is what makes account-level questions possible, and it is the basis of the at-risk accounts table on the Customer Health view.

**One caveat worth stating plainly.** Joining one sales record to many support tickets multiplies rows. Left unhandled it would inflate any total. Every figure on the dashboard is calculated to count each deal, lead and ticket once, and the totals were checked against the raw files to confirm it. The numbers on screen are correct.

## Why the dashboard is built the way it is

---

Each of the three views answers the questions you listed, and nothing else. Marketing Funnel covers channel performance, lead-to-deal conversion, campaign return and response rates. Sales Pipeline covers outcomes and loss reasons, revenue by segment and product, win rates by region, representative and product, days to close and average deal value. Customer Health covers ticket volume by account, resolution times by priority, product areas generating load, sentiment, and the risk list.

Every view opens with a row of summary cards, the headline numbers, so the state of the business is readable in about five seconds without touching a filter. The detail sits underneath for anyone who wants it.

Three things connect the views. Selecting a segment such as Enterprise on one view applies to all three, so the picture stays consistent as you move between them. Clicking a row in the at-risk table filters the charts around it. And selecting an account carries that account across all three views at once, which is the link between marketing, sales and support you asked for.

Four written notes sit directly on the dashboard, next to the charts they explain. Each states a number, says why it matters and names an action, so the interpretation travels with the chart rather than living in someone's head.

## Asking questions in plain English, and what configuration changed

---

Alongside the dashboard, the team can ask questions in plain English through Quick Chat. Nine questions were tested twice, once with no configuration and again after building what Amazon Quick calls a Topic, a written description of what the business means by each field, including alternative names for the same thing.

**The configuration produced a clear gain on straightforward questions.** Asked for average deal size by company size, the tool first averaged all 499 deals, won and lost together, and reported \$1,589 for Enterprise. The dashboard reports \$2,515, because it counts only deals we actually won. Both were internally consistent, they were simply answering about different sets of deals. After configuration the tool defaults to won deals and returns \$2,515, matching the dashboard exactly. It also stopped confusing "campaign source" with an internal country code, and corrected an arithmetic mistake it had made in its own written summary, understating campaign spend by \$1.16M.

**It also produced one regression worth knowing about.** Questions that need two sources at once, such as ranking accounts by ticket volume alongside their revenue, were answered correctly before configuration and only partly afterwards. The configuration is scoped to the three original sources and does not include the combined dataset, so the tool lost the bridge between support tickets and sales deals. This is fixable, and it is the first improvement I would make next.

## What the data says, and what I recommend

| What we found                                           | The numbers                                                                                                                                                                                                | Recommended action                                                                                                                                                                                                                   |
|---------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Every campaign is underwater on attributed revenue      | \$12.36M spent across six campaigns returned \$1.13M in attributed revenue, an average of -90.3%. NovaEdge Awareness is worst at -97.7%.                                                                   | Do not cut budget on these figures alone. A gap this uniform across all six points at how revenue is credited back to campaigns rather than at six separately failing campaigns. Audit the attribution rules first.                  |
| The channels that convert best get the least volume     | Direct Mail converts at roughly 45% and Partner Referral at roughly 30%. Organic Search sits near 2% and Email near 5%, yet Organic Search carries the largest share of leads.                             | Shift spend toward Direct Mail and Partner Referral, and set a floor on the two low-converting channels. Confirm how much volume each can absorb before moving money.                                                                |
| A high win rate is hiding small deals                   | Central wins 69.9% of its deals at an average of \$1,905 per won deal. West wins 58.8% but averages \$2,559.                                                                                               | Set rep targets on revenue per won deal, not win rate alone. On the current measure Central looks strongest while bringing in less money per deal than West.                                                                         |
| Nearly half of all losses were avoidable                | Of 184 lost deals, No Decision Made (43) and Poor Product Fit (43) account for 86, that is 47% of losses. Neither is about price or a competitor.                                                          | Require a qualification check, budget confirmed, decision maker identified and product fit verified, before a deal moves to Opportunity stage. The average deal runs 67 days, so this is a large amount of recoverable selling time. |
| Urgent support tickets are not treated as urgent        | Resolution time is flat across priority levels. Critical tickets average 56.6 hours and low-priority 59.4 hours. The spread across all six product areas is under two hours.                               | Define priority-based service levels, starting with an eight-hour target for critical tickets. Severity currently has no bearing on how fast an issue gets fixed.                                                                    |
| Eight accounts are at real risk of churning quietly     | Eight accounts combine above-median ticket volume, 25% or more negative sentiment and above-median revenue. Falcon Software is clearest, \$20,912 in won revenue with 41% of its tickets negative in tone. | Run executive account reviews on this list this quarter. The dashboard flags them automatically as data refreshes, so this becomes a standing agenda item.                                                                           |
| Our largest customer is also our largest support burden | YieldMax Software files 334 support tickets, nearly double the next account, and holds the highest won revenue at \$40,722.                                                                                | Assign dedicated support coverage. The account carries the most revenue and the most operational load at once, which is the profile that churns without warning.                                                                     |

## Where the AI agreed with the dashboard, and where it did not

On single, clearly defined measures the two agree exactly. Won revenue of \$707,201, average won deal of \$2,245, campaign spend of \$12,359,497, attributed revenue of \$1,127,223 and average resolution of 58.2 hours all came back identical from both. Where a question maps to one number from one source, the plain-English tool can be trusted.

The disagreements were all about definitions rather than arithmetic, and they fall into three groups. The first is which records to count, won deals against all deals, which was the \$1,589 against \$2,515 case above. The second is words that sound precise but are not defined anywhere in the data. "Conversion rate" is the clearest example, the tool still reads it as a count of customers who responded rather than the share of leads that became deals, so its channel ranking does not match the dashboard. That one needs a definition agreed by the business, not a technical fix. The third is questions that span two sources, described in the previous section.

**The practical guidance for the team is short.** Use the dashboard for the numbers reviewed every Monday, where the definition has to be the same every week and people need to filter and drill in. Use plain-English questions for the one-off nobody anticipated, and for exploring an idea before deciding whether it deserves a permanent chart. When a plain-English answer is going to inform a decision, read the line showing how the question was interpreted before trusting the number. Every incorrect answer in the testing announced itself on that line first.

## What I would do next

- Agree a definition of conversion rate.** It is the one term where the dashboard and the plain-English tool genuinely disagree, and it needs a business decision rather than a technical one.
- Extend the plain-English configuration to the combined dataset** so account-level questions spanning marketing, sales and support work as well there as they do on the dashboard.
- Review campaign attribution** before acting on the return figures. A -90.3% average across all six campaigns is more likely to be a measurement problem than six independent failures.
- Put the eight at-risk accounts on a standing quarterly review.** The list refreshes itself as the data refreshes, so it does not need rebuilding each time.

Terms used in this report: SPICE is the in-memory store that keeps the dashboard fast. A Topic is a written description of the data that helps the plain-English tool interpret business terms correctly. Attributed revenue is revenue credited back to the campaign that generated the lead. Sentiment is the tone recorded on

*a support ticket, positive, neutral or negative. Win rate is won deals as a share of all closed deals.*